



AI CODING SPEND · PROTOTYPING

drafts & spikes	\$0.00
refactors	\$0.00
throwaway tests	\$0.00


TOTAL

\$0



OPEN SOURCE · MIT · 28.3K ★

Prototype on free tokens.
Ship on paid ones.





THE PROBLEM

Exploration is the cost nobody budgets for.


Most of what a coding agent writes in the first week of a project gets thrown away. Spikes, dead ends, three versions of the same component.

Paying frontier rates for that, or stopping when a usage limit lands mid-task, is a choice. There's now a free one.

v1
Checkout.tsx
× THROWN AWAY

v2
Checkout.tsx
× THROWN AWAY

v3
Checkout.tsx
✓ SHIPPED





THE TOOL

FreeLLMAPI

An open-source router you run locally. It puts every free LLM tier behind one /v1 endpoint that speaks both OpenAI and Anthropic.

7.4B

TOKENS / MONTH


34

PROVIDERS


635

ENDPOINTS


github.com/tashfeenahmed/freellmapi




Groq/




Gemini/



Mistral/



NVIDIA/






HOW IT ROUTES

Four rules, doing the work for you.

- 1 Ranks.** Every model gets a live score for speed, capability and reliability.
- 2 Fails over.** A rate limit sends the call to the next model, with cooldowns and key rotation.
- 3 Stays put.** A conversation keeps one model for 30 minutes, so it isn't hopping every turn.
- 4 Hands off.** When it does switch, a compact note carries the thread across.





SETUP

Three commands.

terminal

1. install, then open localhost:3001


```
$ curl -fsSL https://freellmapi.co/install.sh | bash
```

2. add your free provider keys

in the dashboard

3. point Claude Code at it

```
$ npx freellmapi setup-claude \  
  --url http://localhost:3001 \  
  --api-key <unified-key>
```





WHERE IT FITS

Good for this. Not for that.


USE IT FOR

- ✓ Prototypes and spikes
- ✓ Side projects
- ✓ Learning an agent workflow
- ✓ Bulk drafts you'll rewrite

NOT FOR

- ✗ Production
- ✗ Client or private data: calls go to the free providers
- ✗ Work that needs a frontier model

Its own README: no frontier models, variable latency, no SLA, and weaker late in the day until the caps reset at UTC midnight.







THE RULE WORTH KEEPING

Free for exploring. Paid for shipping.


Send the throwaway work to free tiers. Keep the paid model for the final pass and for anything a customer will touch.





Follow Yar Malik

for Claude Code setups that actually hold up.



comment

"FREE"

and I'll send the repo and the setup.

